



# Time Series Analysis

---

# Recap

In the last section, we expanded our framework to include state space models

- There is a hidden or latent process,  $x_t$ , called the state process.  
Assumed to be a Markov process
- There are observations,  $y_t$ , that are generated by the state  
but are otherwise independent
- Linear state space models are a superset of our SARIMAX family of models  
We can rewrite them in state space form
- Dynamic Linear regression  
The weighting coefficients are allowed to change over time
- Kalman Filters  
Enables optimal estimation of state in the presence of noisy measurements

In this section, we'll explore Hidden Markov Models

- Markov Chains
- Viterbi, etc.

# State Space Models

When we introduced state space models, we said that they followed two main principles:

1. There is a *hidden* or *latent* process,  $x_t$ , called the state process.
  - The state process is assumed to be a Markov process
2. The observations,  $y_t$ , are independent given the states  $x_t$ 
  - Any relationship between observations is dictated by the originating states

We've been focused on linear Gaussian state space models, but this is just one (albeit popular) case.

Another well-known case is when the hidden state,  $x_t$ , follows a discrete-valued Markov chain.


- The distribution of the possible observations at some time,  $t$ , is dictated by the current state at that time.
- The current state is one of  $m$  possible states, and the states evolve according to a Markov Chain over time.
- When the underlying state is unknown (we only have the observations), it is termed a hidden Markov model (HMM)

# Markov Process & Chains

A Markov process is a stochastic process whose future value is only dependent on its current state - not its past.


$$p(x_t | x_{t-1}, x_{t-2}, \dots, x_{t-\infty}) = p(x_t | x_{t-1})$$

- A discrete Markov chain is a specific type of Markoff process, which has a countable number of states.
- Time can also be continuous or discrete (but we'll focus on discrete time)
- In a [hidden Markov model](#), the state is hidden, and we only see a corresponding observation




# Markov Process & Chains

- Hidden States: Phonemes of speech
- Observations: Acoustic Spectra



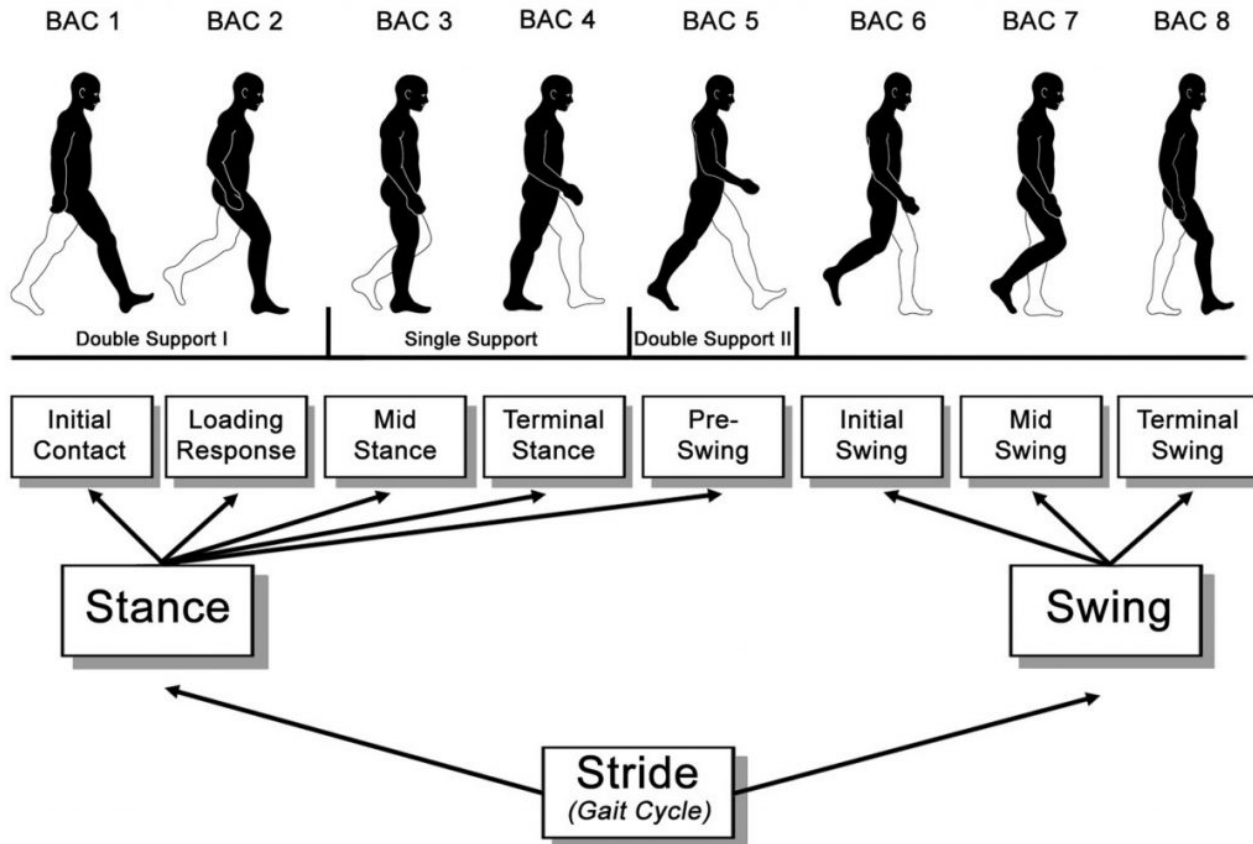
- Hidden States: Stages of sleep
- Observations: EEG signals



# Markov Process & Chains

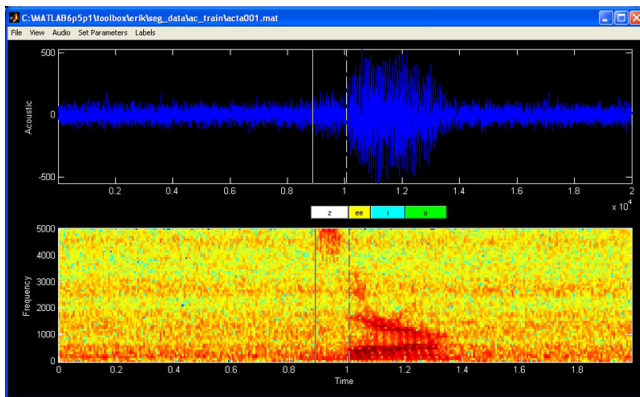
Examples:

- Hidden States: Gait Phase
- Observations: IMU or Video data




# Markov Process & Chains

E. J. Scheme, B. Hudgins and P. A. Parker, "Myoelectric Signal Classification for Phoneme-Based Speech Recognition," in *IEEE Transactions on Biomedical Engineering*, vol. 54, no. 4, pp. 694-699, April 2007, doi: 10.1109/TBME.2006.889175.



Word	Phonemes			
zero:	/ Z /	/ i /	/ r /	/ o /
one:	/ w /	/ A /	/ n /	
two:	/ t /	/ u /		
three:	/ T /	/ r /	/ i /	
four:	/ f /	/ o /	/ r /	
five:	/ f /	/ a l /	/ v /	
six:	/ s /	/ I /	/ x /	
seven:	/ s /	/ E /	/ v /	/ n /
eight:	/ e /	/ t /		
nine:	/ n /	/ a l /	/ n /	



# HMMs vs Kalman

In [Kalman state space models](#), we assumed that:

- The unobserved state and the observations were Gaussian
- The model evolves continuously according to linear dynamics

In [HMMs](#), we generally assume that:

- The hidden state is one of a set number of classes (discrete)
- The state/class changes according to a discrete Markov chain
- The observations may be discrete or continuous

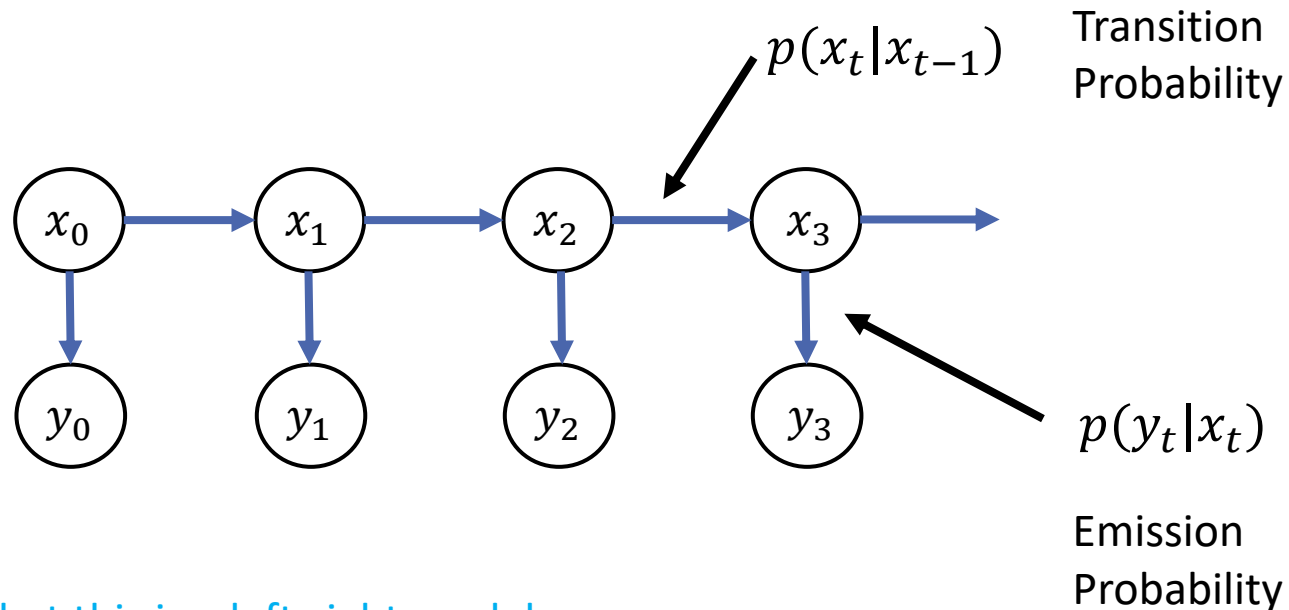


# Markov Process & Chains

In a discrete Markov chain, the observations are deterministic for each given state

- The state IS the observable event, so it is not hidden

A **discrete hidden Markov model** is an extension, where the observation is a probabilistic function of the state



Note that this is a left-right model


# Markov Process & Chains

Below, is a 3-state **Markov Chain**, with transition probabilities:

$$a_{ij} = p(x_j|x_i)$$

$$a_{ij} \geq 0$$

$$\sum_{j=0}^{N-1} a_{ij} = 1$$



For an N-state chain, we get the State Transition Matrix, A

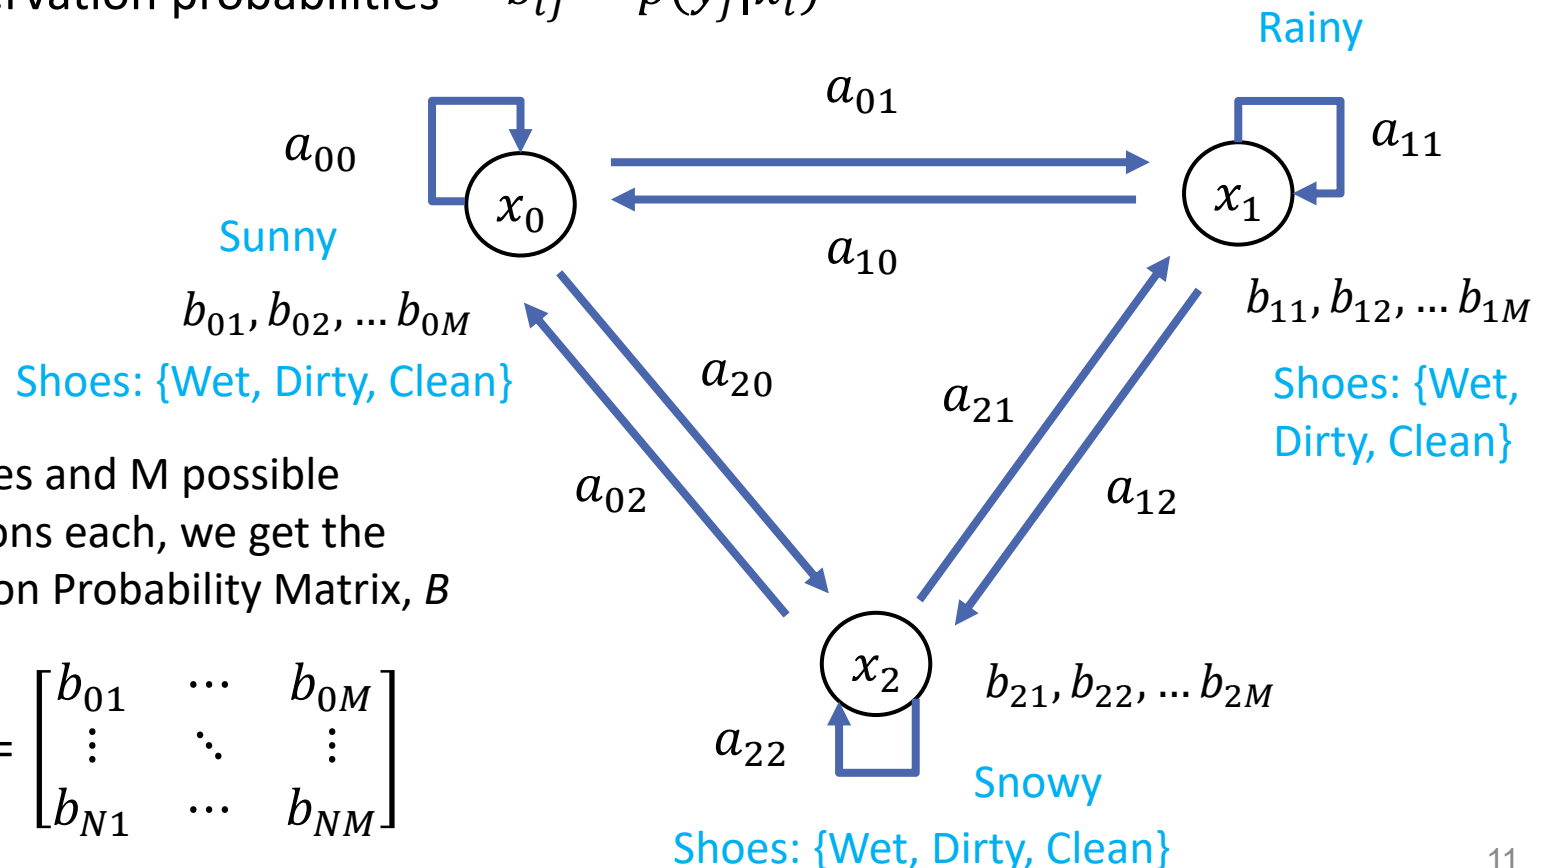
$$A = \begin{bmatrix} a_{00} & \cdots & a_{0N} \\ \vdots & \ddots & \vdots \\ a_{N0} & \cdots & a_{NN} \end{bmatrix}$$

# Hidden Markov Model

For a similar 3-state [Hidden Markov Model](#), we have transition probabilities:

$$a_{ij} = p(x_j|x_i) \quad a_{ij} \geq 0 \quad \sum_{j=0}^{N-1} a_{ij} = 1$$

And observation probabilities  $b_{ij} = p(y_j|x_i)$



For N states and M possible observations each, we get the Observation Probability Matrix,  $B$

$$B = \begin{bmatrix} b_{01} & \cdots & b_{0M} \\ \vdots & \ddots & \vdots \\ b_{N1} & \cdots & b_{NM} \end{bmatrix}$$

# HMM Example


Let's look at a series of **coin tosses** as a simple example:

- Someone is flipping one or more coins (you can't see them)
- The observation sequence is the corresponding series of heads and tails

Example observations:  $O = \{o_1, o_2, o_3 \dots o_T\} = \{H H T \dots T\}$

We could establish several different HMM models for the problem

- M1: Assume one coin




Observation,  $O = \{H H T T H T H H T T H\}$   
State Sequence,  $Q = \{1 1 2 2 1 2 1 1 2 2 1\}$

# HMM Example

We could establish several different HMM models for the problem

- M2: Assume two coins



Observation,  $O=\{H \ H \ T \ T \ H \ T \ H \ H \ T \ T \ H\}$   
State Sequence,  $Q=\{2 \ 1 \ 1 \ 2 \ 2 \ 2 \ 1 \ 2 \ 2 \ 1 \ 2\}$

- In this case, we don't know if the coins are biased (e.g.  $p(H) \neq 0.5$ ) AND we don't know if they are biased equally
- There is a non-zero probability that the same flip result (H or T) can come from either coin

# HMM Example


We could establish several different HMM models for the problem

- M3: We could similarly assume that we have three coins, but this would look just like our previous 3-state example
- We might get an example sequence of:

Observation,  $O=\{H \ H \ T \ T \ H \ T \ H \ H \ T \ T \ H\}$   
State Sequence,  $Q=\{3 \ 1 \ 2 \ 3 \ 3 \ 1 \ 1 \ 2 \ 3 \ 1 \ 3\}$

We can explain a sequence of observations with a variety of different models

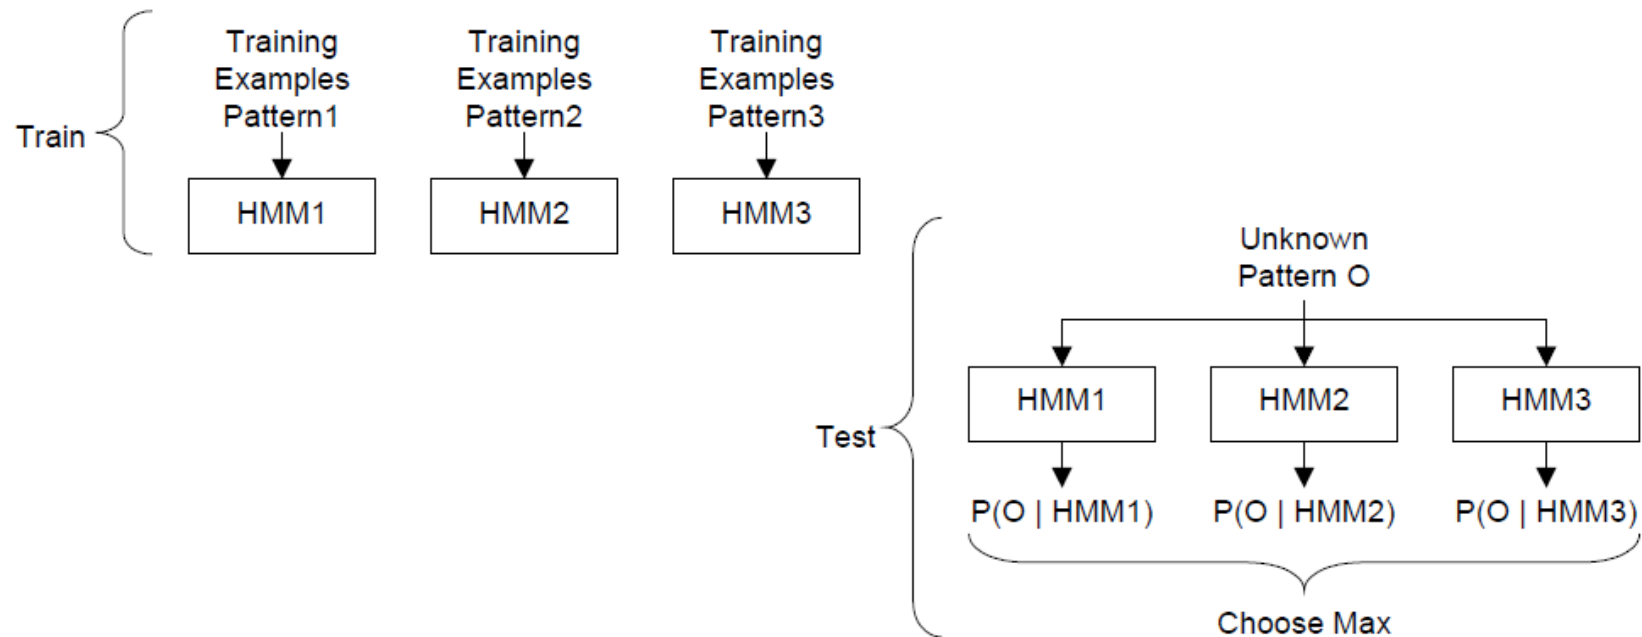
- In the 2-state example, there were  $2^T$  possible state sequences, whereas
- In the 3-state case, there were  $3^T$  possible state sequences...
- There are an infinite number of HMMs that could explain that particular sequence
- As in all ML, there is a capacity tradeoff, so choice is important



# Uses of HMMs

There are two main uses of HMM, in practice

- **Modeling**: match an HMM to an observation sequence to gain a better understanding of the process
- **Pattern Classification**: Leverage the temporal information to classify a sequence (e.g. Sign language gesture recognition)



# Hidden Markov Models

HMMs are defined by a set of 3 main components

For a set of  $N$  hidden states,  $S = \{S_1, S_2 \dots S_N\}$

and a set of  $M$  possible observations,  $O = \{O_1, O_2, \dots O_M\}$

a. State transition probability matrix,  $A$  of size  $[N, N]$ , composed of

$$a_{ij} = P(q_{t+1} = S_j | q_t = S_i) \quad (q_t \text{ is the hidden state at time } t)$$

b. State emission probabilities,  $B$  of size  $[N, M]$ , composed of

$$b_j(k) = P(o_t = O_k | q_t = S_j) \quad (o_t \text{ is the observation at time } t)$$

c. Initial distribution vector,  $\pi$  of size  $[N]$ , where

$$\pi_i = P(q_1 = S_i)$$



# Hidden Markov Models

From a list of 3000 words from the Oxford Advanced Learner's Dictionary, find the

- Probability of having any word start with a given letter
- Probability of a transition from any one letter to the next

	a	b	c	d	e	f	g	h
<start>	7.4	4.8	10.1	5.4	5.5	4.7	2.6	3.1
a	0.0	3.1	6.1	4.9	0.1	1.1	3.3	0.2
b	13.7	0.0	0.0	0.0	17.4	0.0	0.0	0.0
c	11.4	0.0	2.3	0.0	15.8	0.0	0.0	10.7
d	6.7	0.0	0.0	2.4	39.2	0.5	1.7	0.2
e	8.9	0.4	5.5	3.7	3.6	2.1	1.9	0.3
f	12.9	0.0	0.0	0.0	15.9	9.1	0.0	0.0
g	11.2	0.0	0.0	0.0	26.8	0.0	1.6	13.1
h	15.5	0.5	0.0	0.5	26.7	0.0	0.0	0.0

1<sup>st</sup> row: Initial distribution vector

- 10.1% chance that a word starts with the letter C.

All other rows: transition probability matrix,  $A$

- 39.2% chance that the letter  $D$  is followed by an  $E$ .

Question: What use case might lead to non-deterministic state emission probabilities?

# Hidden Markov Models

From a list of 3000 words from the Oxford Advanced Learner's Dictionary, find the

- Probability of having any word start with a given letter
- Probability of a transition from any one letter to the next

Here are some **state emission probabilities, B**, for a case where you type the right key 70% of the time.

They assume that you hit one of the surrounding keys the rest of the time.

	a	b	c	d	e	f	g	h
a	70.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
b	0.0	70.0	0.0	0.0	0.0	0.0	7.5	0.0
c	0.0	0.0	70.0	7.5	0.0	0.0	0.0	0.0
d	0.0	0.0	5.0	70.0	5.0	5.0	0.0	0.0
e	0.0	0.0	0.0	7.5	70.0	0.0	0.0	0.0
f	0.0	0.0	5.0	5.0	0.0	70.0	5.0	0.0
g	0.0	5.0	0.0	0.0	0.0	5.0	70.0	0.0
h	0.0	5.0	0.0	0.0	0.0	0.0	5.0	70.0

# Hidden Markov Models

The math associated with HMMs can then be broken into 3 main problems

- **Evaluation**: Given an HMM, say  $\lambda$ , and a sequence of observations, say  $O$ , find the probability of that set of observations being explained (generated) by that model (**Forward, Backward algorithms**)

$$P(O|\lambda)$$

- **Decoding**: Given an HMM,  $\lambda$ , and a sequence of observations,  $O$ , find the sequence of hidden states,  $Q$  that maximizes (**Viterbi algorithm**)

$$P(O, Q|\lambda)$$

- **Learning**: Given an unknown HMM,  $\lambda$ , and a sequence of observations,  $O$ , find the parameters of  $\lambda(A, B, \pi)$  that maximize: (**Baum-Welch, Forward-Backward algorithm**)

$$P(O|\lambda(A, B, \pi))$$

# Evaluation

**Evaluation:** Given an HMM, say  $\lambda$ , and a sequence of observations, say  $O$ , find the probability of that set of observations being explained (generated) by that model:

$$P(O|\lambda)$$

Computing this term allows us to determine how well an HMM matches a given observation sequence.

# Decoding

**Decoding:** Given an HMM,  $\lambda$ , and a sequence of observations,  $O$ , find the sequence of hidden states,  $Q$  that maximizes:

$$P(O, Q | \lambda)$$

Determine the state sequence with the highest probability for a given observation sequence and HMM.

$$\operatorname{argmax}\{P(Q | O, \lambda)\}$$

This optimal state sequence,  $Q_{opt}$ , is not necessarily the correct state sequence (the true sequence that generated the observation), but it gives some indication about the average statistics of the model and individual states

# Learning

**Learning:** Given an unknown HMM,  $\lambda$ , and a sequence of observations,  $O$ , find the parameters of  $\lambda(A, B, \pi)$  that maximize:

$$P(O|\lambda(A, B, \pi))$$

The solution of this problem lets us to *train* an HMM based on a set of observed training data.

# Evaluation: Direct Computation

There are different ways to execute these three different problems when working with HMMs. Let's first look at the **evaluation** problem,  $P(O|\lambda)$

For a given HMM, we can calculate the probability of a particular observation and state sequence as

$$P(O, Q|\lambda) = \pi_1 b_{q_1}(o_1) a_{q_1 q_2} b_{q_2}(o_2) a_{q_2 q_3} b_{q_3}(o_3) \dots a_{q_{T-1} q_T} b_{q_T}(o_T)$$

$\pi_1$  is the probability of starting in state  $q_1$

$b_{q_1}(o_1)$  is the probability of observing  $o_1$ , given that we are in state  $q_1$

$a_{q_1 q_2}$  is the probability of a state transition from  $q_1$  to  $q_2$

$b_{q_2}(o_2)$  is then the probability of observing  $o_2$ , given that we are in state  $q_2$

# Evaluation: Direct Computation

$$P(O, Q|\lambda) = \pi_1 b_{q_1}(o_1) a_{q_1 q_2} b_{q_2}(o_2) a_{q_2 q_3} b_{q_3}(o_3) \dots a_{q_{T-1} q_T} b_{q_T}(o_T)$$

This is the probability for one specific set of observations AND state sequence, though.

- To determine the probability of a set of observations for a given HMM, we must use the equation determine above for *every possible state sequence* and sum them all.

$$P(O|\lambda) = \sum_{all Q} P(O, Q|\lambda)$$

- Each  $P(O, Q|\lambda)$  requires the multiplication of  $2T$  terms (where  $T$  is the number of observations)
- For an  $N$  state HMM, there are  $N^T$  of the terms because at every  $t$  there are  $N$  possible states that can be reached.
- Even for small values of  $N$  and  $T$ , computing  $P(O|\lambda)$  becomes impractical  
(e.g.  $N = 5$ ,  $T = 100$ ,  $2T \cdot N^T = 1.57 \times 10^{72}$ )
- Clearly, direct computation of  $P(O|\lambda)$  is unfeasible and a more efficient procedure is required.



# Evaluation: Forward Procedure

Instead, we can define the forward variable  $\alpha_t(i)$

$$\alpha_t(i) = P(o_1 o_2 \dots o_t, q_t = i | \lambda)$$

as the probability of a partial observation (an observation sequence,  $O$ , up to time  $t$ ), AND being in state  $i$  at time  $t$ .

This is known as the **forward procedure** because we start with a partial observation of  $o_1$  and move forward, solving for  $\alpha_t(i)$  iteratively until time  $T$  so that the entire observation sequence is taken into account.

## 1. Initialization

$$\alpha_1(i) = \pi_i b_i(o_1) \quad 1 \leq i \leq N$$

The forward variable  $\alpha_1(i)$  is the probability of observing  $o_1$  in state  $i$ .

- the probability of starting in state  $i$  multiplied by the probability of observing  $o_1$  while in state  $i$ .

# Evaluation: Forward Procedure

## 2. Induction


$$\alpha_{t+1}(j) = \left( \sum_{i=1}^N \alpha_t(i) a_{ij} \right) b_j(o_{t+1}) \quad 1 \leq t \leq T - 1 \quad 1 \leq j \leq N$$

In the induction step, the probability of the partial observation sequence up to time  $t$  is assumed to be known.

We can then determine the probability of being in state  $j$  at the next time step ( $t + 1$ ) by summing the  $\alpha_t(i)a_{ij}$  terms, the probabilities of being in state  $i$  and having a state transition into state  $j$

The  $b_j(o_{t+1})$  term accounts for the probability of actually observing  $o_{t+1}$  given that we are to be in state  $j$

# Evaluation: Forward Procedure



# Evaluation: Forward Procedure

## 3. Termination

$$P(O|\lambda) = \sum_{i=1}^N \alpha_T(i)$$

At this point, the forward variable at time  $T$  is no longer a partial observation because it takes into account the entire observation sequence  $O$ .

- The term  $\alpha_T(i) = P(O, q_T = i|\lambda)$  is the probability of observing  $O$  and ending in state  $i$ .
- $P(O|\lambda)$  is then just the sum of  $\alpha_T(i)$  for all possible ending states.

The calculation of the forward algorithm is on the order of  $N^2T$ .

(e.g.  $N = 5, T = 100, N^2 \cdot T = 100 \cdot 5^2$  instead of  $2T \cdot N^T = 1.57 \times 10^{72}$ )

# Evaluation: Forward Procedure

## 1. Initialization

- Determine the forward variable the probability of observing  $o_1$  in state  $i$ .

$$\alpha_1(i) = \pi_i b_i(o_1)$$

## 2. Induction

- We determine the probability of going to state  $j$  at the next time step

$$\alpha_{t+1}(j) = \left( \sum_{i=1}^N \alpha_t(i) a_{ij} \right) b_j(o_{t+1})$$

## 3. Termination

- We sum the probabilities of observing the sequence  $O$  and ending in state  $i$

$$P(O|\lambda) = \sum_{i=1}^N \alpha_T(i)$$

Note: In the time-series sense, this operation can be viewed as *filtering*

# Evaluation: Backward Procedure

The **backward procedure** is like the forward procedure except that we instead start at time  $T$  and work our way backwards in time, solving iteratively for the backward variable  $\beta_t(i)$ .

$$\beta_t(i) = P(o_{t+1}o_{t+2} \dots o_T | q_t = i, \lambda)$$

This *backward variable* is the probability of the partial observation (observation sequence from time  $t + 1$  to time  $T$ ) given that at time  $t$  we are in state  $i$ .

- In the forward variable, the probability was a joint probability between the partial observation and the state  $q_t$

$$\alpha_t(i) = P(o_1o_2 \dots o_t, q_t = i | \lambda)$$

- In the backward variable, it is conditional probability between the partial observation and the state  $q_t$

# Evaluation: Backward Procedure

## 1. Initialization

$$\beta_T(i) = 1 \quad 1 \leq i \leq N$$

We're starting at the end, so we can't define the partial observation sequence from  $o_{T+1}$  (beyond our sequence) to  $o_T$ . So, for the procedure to work we set  $\beta_T(i)$  equal to 1.

# Evaluation: Backward Procedure

## 2. Induction

$$\beta_t(i) = \sum_{j=1}^N a_{ij} b_j(o_{t+1}) \beta_{t+1}(j) \quad T-1 \geq t \geq 1 \quad 1 \leq i \leq N$$

In the backward induction step, we know the backward variable from time  $t + 1$  to time  $T$


$\beta_{t+1}(j)$  accounts for everything that happens after  $t + 1$ , given that we are in state  $j$  at time  $t + 1$

For a given state  $i$  at time  $t$ , the induction step must consider transitioning into every state  $j = 1, 2, \dots, N$

Because we know we'll be in  $\beta_{t+1}(j)$ , we need the single observation probability  $b_j(o_{t+1})$ , and the transition probabilities  $a_{ij}$



# Evaluation: Backward Procedure



# Evaluation: Backward Procedure

## 3. Termination

$$P(O|\lambda) = \sum_{i=1}^N \pi_i b_i(o_1) \beta_1(i)$$

Finally,  $\beta_1(i)$  is the probability of the partial observation  $\{o_2 o_3 \dots o_T\}$  given that the initial state is state  $i$ .

The  $\pi_i$  removes the condition on the probability that the initial state is given to be state  $i$ .

The  $b_i(o_1)$  is the probability of seeing the observation  $o_1$  at the initial state  $i$ .

To compute  $P(O|\lambda)$  we need to sum to account for all possible initial states.

Note: In the time-series sense, this operation can be viewed as *smoothing*

# Evaluation: Backward Procedure

## 1. Initialization

- We require that  $\beta_T(i) = 1 \quad 1 \leq i \leq N$

## 2. Induction

- We determine the probability of going to  $\beta_t(i)$  given  $\beta_{t+1}(j)$

$$\beta_t(i) = \sum_{j=1}^N a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)$$

## 3. Termination

- We sum the probabilities of the partial observations  $\{o_2 o_3 \dots o_T\}$  given the possible initial states  $i$

$$P(O|\lambda) = \sum_{i=1}^N \pi_i b_i(o_1) \beta_1(i)$$

# Decoding: Viterbi Algorithm

Remember, in the **decoding** problem, we need to determine the optimal state sequence for a given set of observations.

- The criteria we used to define 'optimal' was the state sequence that maximized  $P(Q|O, \lambda)$ .

The formal solution for this is called the **Viterbi algorithm**.

We define  $\delta_t(i)$  to be the highest probability for a partial state sequence  $\{q_1, q_2, \dots, q_t\}$  ending at state  $q_t = i$  for a partial observation sequence  $\{o_1, o_2, \dots, o_t\}$ .

$$\delta_t(i) = \max_{q_1, q_2, \dots, q_{t-1}} [P(q_1, q_2, \dots, q_{t-1}, q_t = i, o_1, o_2, \dots, o_t | \lambda)]$$

The calculation of the Viterbi algorithm is accomplished using **dynamic programming**.

# Decoding: Viterbi Algorithm

We can solve for  $\delta_t(i)$  inductively using

$$\delta_{t+1}(j) = \max_i [\delta_t(i) a_{ij}] \cdot b_j(o_{t+1})$$

The probability  $\delta_t(i)$  is the probability of the optimal path for each state up to time  $t$ .

To find the optimal path to state  $j$  at time  $t + 1$ , we calculate, for all  $i$ , the probability of moving from state  $i$  (given that it was reached optimally at time  $t$ ) to state  $j$ , and take the path with maximum probability.

The term  $b_j(o_{t+1})$  is required by the definition of  $\delta_t(i)$  because it takes into account the probability of observing  $o_{t+1}$  at time  $t$ .

If we track the optimal state transition at each induction step, then we get the optimal state sequence  $Q_{opt} = \{q_{opt(1)}, q_{opt(2)}, \dots, q_{opt(N)}\}$ .

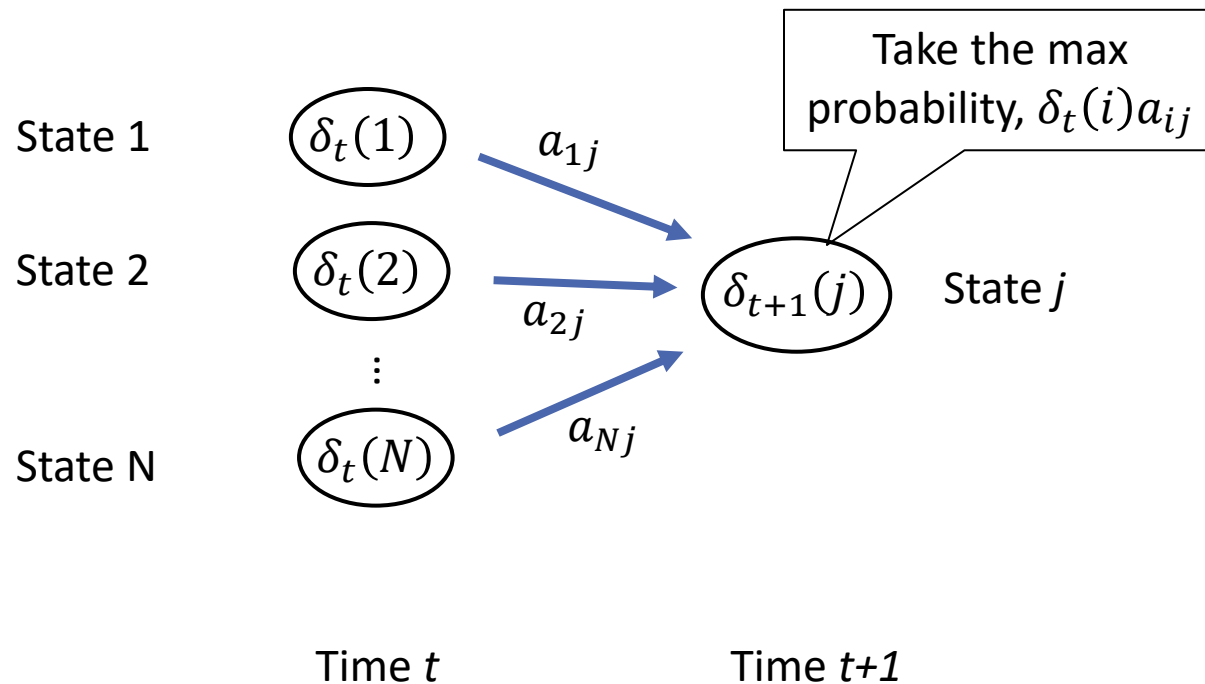
- The array  $\psi_1(j)$  is used to keep track of the optimal state transition argument

# Decoding: Viterbi Algorithm

We can solve for  $\delta_t(i)$  inductively using

$$\delta_{t+1}(j) = \max_i [\delta_t(i) a_{ij}] \cdot b_j(o_{t+1})$$

The probability  $\delta_t(i)$  is the probability of the optimal path for each state up to time  $t$ .



# Decoding: Viterbi Algorithm

## 1. Initialization

$$\delta_1(i) = \pi_i b_i(o_1) \quad 1 \leq i \leq N$$

$$\psi_1(i) = 0 \quad 1 \leq i \leq N$$

For time  $t = 1$ ,  $\delta_1(i)$  is just the probability of the initial state  $i$  multiplied by the probability of observing  $o_1$  when in state  $i$ .

The array  $\psi_1(i)$  is set arbitrarily to zero at time  $t = 1$ , but is never be actually used in the algorithm

- there was no state to transition from prior to this

# Decoding: Viterbi Algorithm

## 2. Recursion

$$\delta_{t+1}(j) = \max_i [\delta_t(i) a_{ij}] \cdot b_j(o_{t+1})$$

$$\psi_t(i) = \arg \left[ \max_i [\delta_t(i) a_{ij}] \right]$$

$$1 \leq j \leq N \quad 2 \leq t \leq T$$

The probability  $\delta_t(i)$  is determined recursively as we just walked through.

The argument that indicates the optimal state transition is stored for each state at each time  $t$ .



# Decoding: Viterbi Algorithm

## 3. Termination

$$P_{opt} = \max_i [\delta_T(i)]$$

$$q_{opt}(T) = \arg \left[ \max_i [\delta_T(i)] \right]$$

$P_{opt}$  is the probability of the optimal state sequence (the sequence with the maximum probability) for the observation sequence

- has an end state of  $q_{opt}(T)$

# Decoding: Viterbi Algorithm

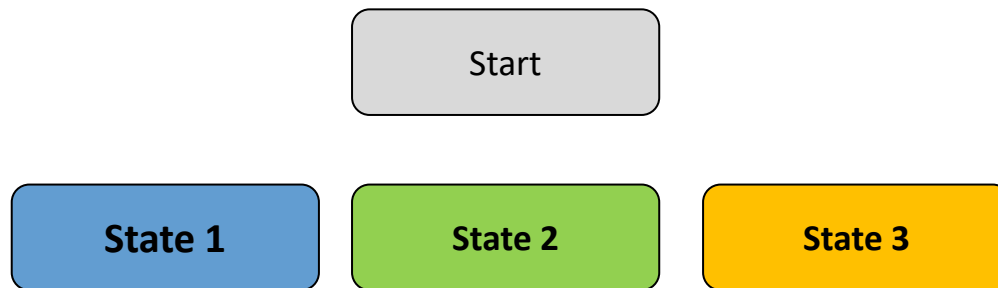
## 4. Path Backtracking

$$q_{opt}(t) = \psi_{t+1}(q_{opt}(t+1)) \quad T-1 \geq t \geq 1$$

Now that we know that the optimal path ends at  $q_{opt}(T)$ , and we know the optimal state transitions, we can trace backwards to find the optimal path.

We start with  $q_{opt}(T)$  and work backwards in time to find the optimal sequence of states,  $Q_{opt}$ .

# Decoding: Viterbi Algorithm



$$\pi = \begin{bmatrix} 1/6 \\ 3/6 \\ 2/6 \end{bmatrix} \quad A = \begin{bmatrix} 4/6 & 1/6 & 1/6 \\ 2/6 & 2/6 & 2/6 \\ 1/6 & 1/6 & 4/6 \end{bmatrix}$$

$$B = \begin{bmatrix} 3/6 & 3/6 \\ 5/6 & 1/6 \\ 2/6 & 4/6 \end{bmatrix}$$

$$O = [1 \quad 2 \quad 1 \quad 1 \quad 2]$$

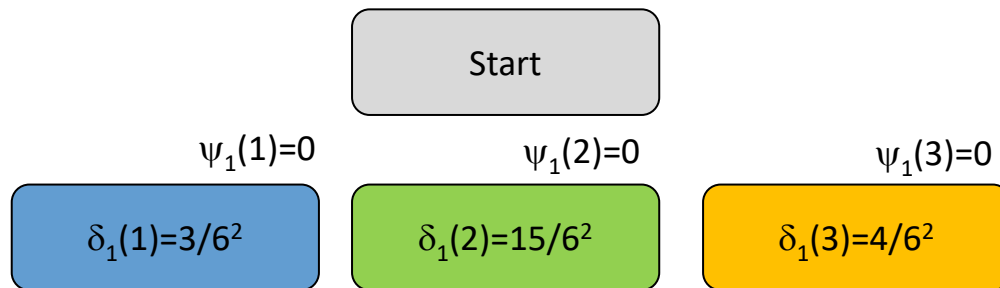
$$Q_{opt} = ?$$

## 1. Initialization

$$\delta_1(i) = \pi_i b_i(o_1)$$

$$\psi_1(i) = 0$$

# Decoding: Viterbi Algorithm



$$\pi = \begin{bmatrix} 1/6 \\ 3/6 \\ 2/6 \end{bmatrix} \quad A = \begin{bmatrix} 4/6 & 1/6 & 1/6 \\ 2/6 & 2/6 & 2/6 \\ 1/6 & 1/6 & 4/6 \end{bmatrix}$$

$$B = \begin{bmatrix} 3/6 & 3/6 \\ 5/6 & 1/6 \\ 2/6 & 4/6 \end{bmatrix}$$

$$O = [1 \quad 2 \quad 1 \quad 1 \quad 2]$$

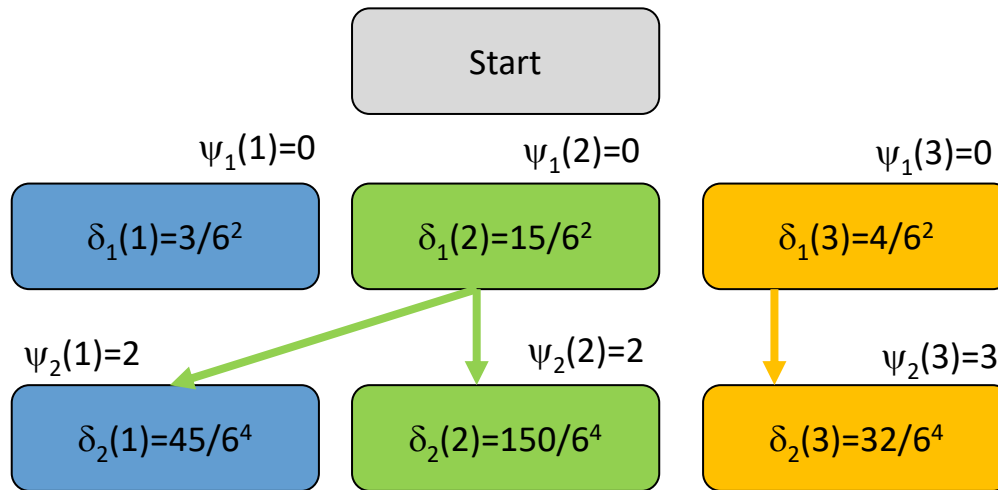
$$Q_{opt} = ?$$

## 1. Initialization

$$\delta_1(i) = \pi_i b_i(o_1)$$

$$\psi_1(i) = 0$$

# Decoding: Viterbi Algorithm



$$\pi = \begin{bmatrix} 1/6 \\ 3/6 \\ 2/6 \end{bmatrix} \quad A = \begin{bmatrix} 4/6 & 1/6 & 1/6 \\ 2/6 & 2/6 & 2/6 \\ 1/6 & 1/6 & 4/6 \end{bmatrix}$$

$$B = \begin{bmatrix} 3/6 & 3/6 \\ 5/6 & 1/6 \\ 2/6 & 4/6 \end{bmatrix}$$

$$O = [1 \quad 2 \quad 1 \quad 1 \quad 2]$$

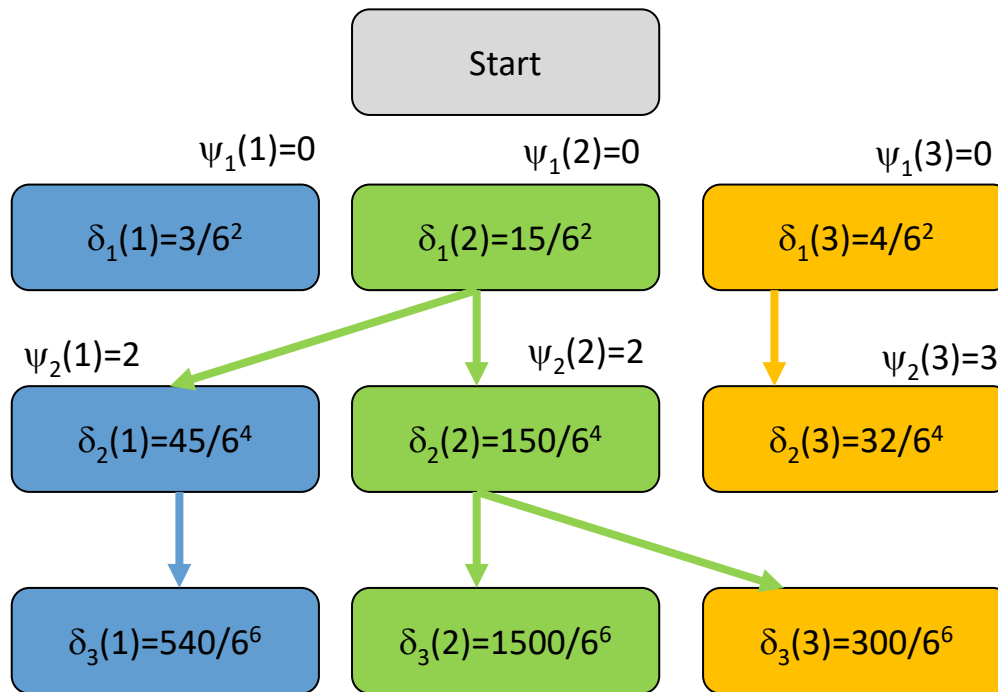
$$Q_{opt} = ?$$

## 2. Recursion

$$\delta_{t+1}(j) = \max_i [\delta_t(i) a_{ij}] \cdot b_j(o_{t+1})$$

$$\psi_t(i) = \arg \left[ \max_i [\delta_t(i) a_{ij}] \right]$$

# Decoding: Viterbi Algorithm



$$\pi = \begin{bmatrix} 1/6 \\ 3/6 \\ 2/6 \end{bmatrix} \quad A = \begin{bmatrix} 4/6 & 1/6 & 1/6 \\ 2/6 & 2/6 & 2/6 \\ 1/6 & 1/6 & 4/6 \end{bmatrix}$$

$$B = \begin{bmatrix} 3/6 & 3/6 \\ 5/6 & 1/6 \\ 2/6 & 4/6 \end{bmatrix}$$

$$O = [1 \quad 2 \quad 1 \quad 1 \quad 2]$$

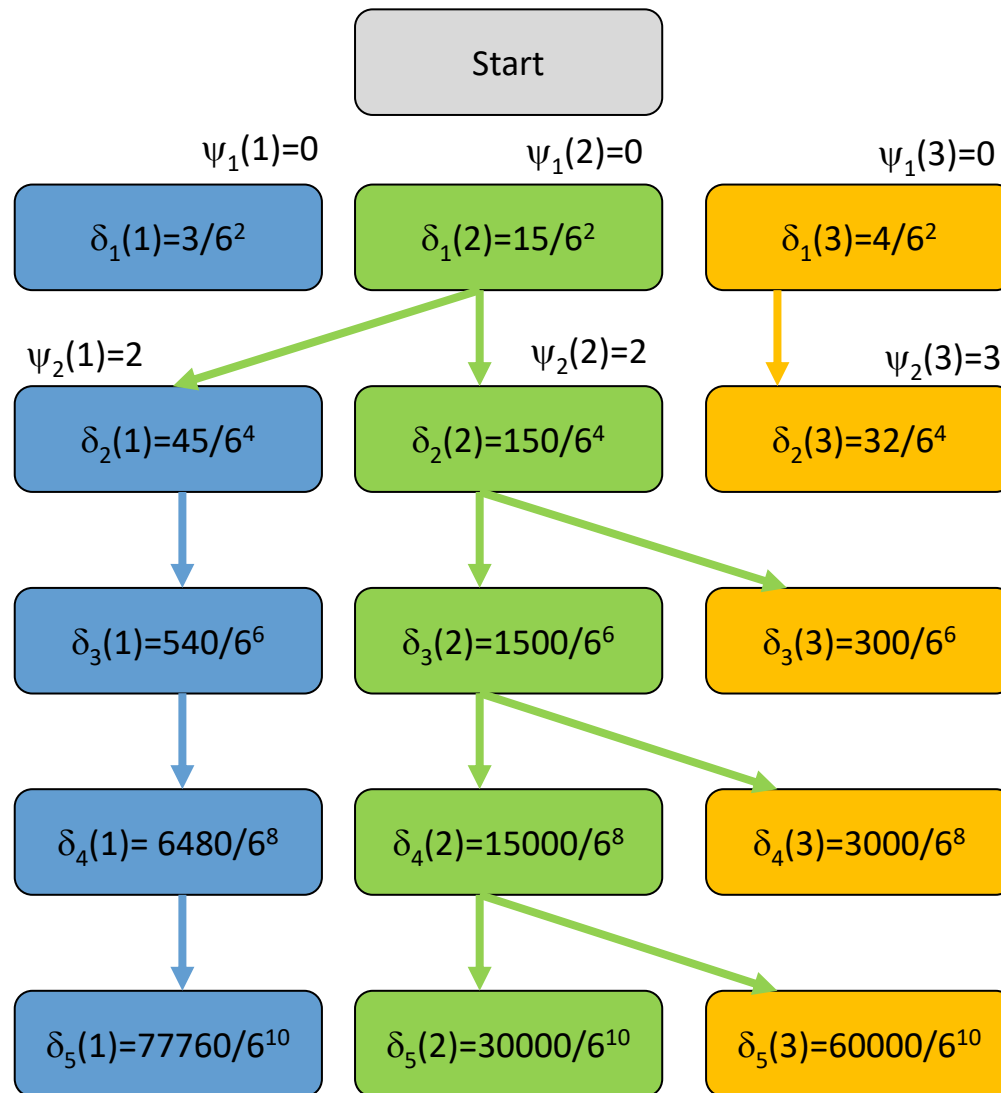
$$Q_{opt} = ?$$

## 2. Recursion

$$\delta_{t+1}(j) = \max_i [\delta_t(i) a_{ij}] \cdot b_j(o_{t+1})$$

$$\psi_t(i) = \arg \left[ \max_i [\delta_t(i) a_{ij}] \right]$$

# Decoding: Viterbi Algorithm



$$\pi = \begin{bmatrix} 1/6 \\ 3/6 \\ 2/6 \end{bmatrix} \quad A = \begin{bmatrix} 4/6 & 1/6 & 1/6 \\ 2/6 & 2/6 & 2/6 \\ 1/6 & 1/6 & 4/6 \end{bmatrix}$$

$$B = \begin{bmatrix} 3/6 & 3/6 \\ 5/6 & 1/6 \\ 2/6 & 4/6 \end{bmatrix}$$

$$O = [1 \quad 2 \quad 1 \quad 1 \quad 2]$$


$$Q_{opt} = ?$$

## 2. Recursion

$$\delta_{t+1}(j) = \max_i [\delta_t(i) a_{ij}] \cdot b_j(o_{t+1})$$

$$\psi_t(i) = \arg \left[ \max_i [\delta_t(i) a_{ij}] \right]$$

# Decoding: Viterbi Algorithm



$$q_{opt}(T)=1, P_{opt} = 77760/6^{10}$$

$$\pi = \begin{bmatrix} 1/6 \\ 3/6 \\ 2/6 \end{bmatrix} \quad A = \begin{bmatrix} 4/6 & 1/6 & 1/6 \\ 2/6 & 2/6 & 2/6 \\ 1/6 & 1/6 & 4/6 \end{bmatrix}$$

$$B = \begin{bmatrix} 3/6 & 3/6 \\ 5/6 & 1/6 \\ 2/6 & 4/6 \end{bmatrix}$$

$$O = [1 \quad 2 \quad 1 \quad 1 \quad 2]$$

$$Q_{opt} = ?$$

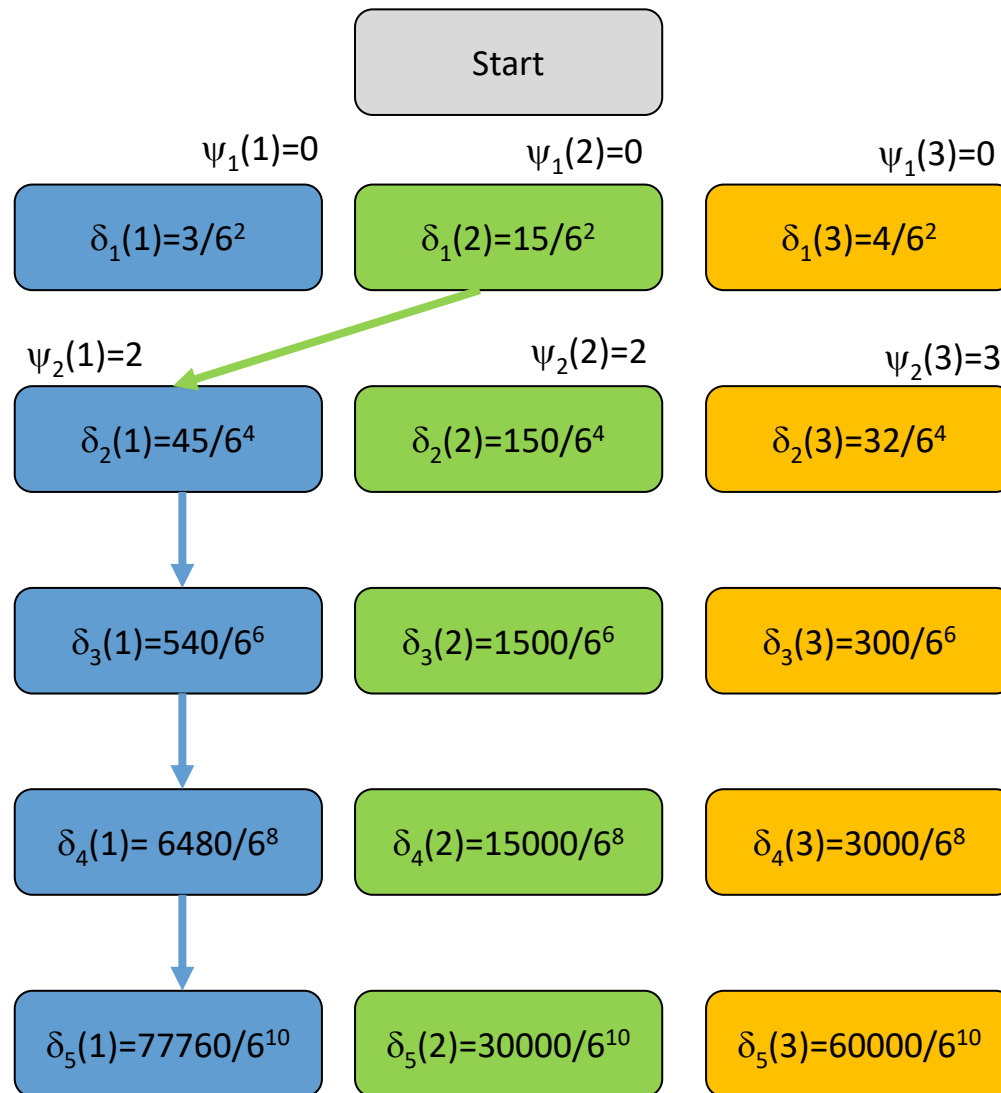
## 3. Termination

$$P_{opt} = \max_i [\delta_T(i)]$$

$$q_{opt}(T) = \arg \left[ \max_i [\delta_T(i)] \right]$$



# Decoding: Viterbi Algorithm



$$q_{opt}(T)=1, P_{opt} = 77760/6^{10}$$

$$\pi = \begin{bmatrix} 1/6 \\ 3/6 \\ 2/6 \end{bmatrix} \quad A = \begin{bmatrix} 4/6 & 1/6 & 1/6 \\ 2/6 & 2/6 & 2/6 \\ 1/6 & 1/6 & 4/6 \end{bmatrix}$$

$$B = \begin{bmatrix} 3/6 & 3/6 \\ 5/6 & 1/6 \\ 2/6 & 4/6 \end{bmatrix}$$

$$O = [1 \quad 2 \quad 1 \quad 1 \quad 2]$$

$$Q_{opt} = ?$$

## 4. Path Backtracking

$$\underline{Q_{opt} = [2 \quad 1 \quad 1 \quad 1 \quad 1]}$$

# Learning (HMM Training)

If you remember in the 3<sup>rd</sup> problem, learning, we needed to find the parameters of an HMM,  $\lambda(A, B, \pi)$ , that maximize the probability of having observed a given sequence,  $O$ :

$$P(O|\lambda(A, B, \pi))$$

The goal here, is therefore to learn, or **train**, the HMM hyperparameters based on a set of training data.

- That is, we need to learn the A (all the  $a_{ij}$ 's), B (all the  $b_j(k)$ ), and  $\pi$
- Here, we'll discuss the popular **Baum-Welch** method (also known as the **forward-backward** method)

# Learning: Baum-Welch Algorithm

The **Baum-Welch** method is a dynamic programming approach and a special case of *expectation maximization* (EM)

- The goal is to iteratively tune the parameters such that we maximize the likelihood that the model “generated” the observations.

## 1. Initialization

- Initialize the  $A, B, \pi$  using prior knowledge if possible (otherwise random)

## 2. Forward Phase

- Compute the forward algorithm  $\alpha_{t+1}(j) = \left( \sum_{i=1}^N \alpha_t(i) a_{ij} \right) b_j(o_{t+1})$

## 3. Backward Phase


- Compute the backward algorithm  $\beta_t(i) = \sum_{j=1}^N a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)$

## 4. Re-Estimation

- Compute the new parameters

# Learning: Baum-Welch Algorithm

First, let's define  $\gamma_t(i)$  as the probability of being in state  $i$  at time  $t$ , given an observation sequence.

$$\gamma_t(i) = P(q_t = i | O, \lambda) = \frac{P(q_t = i | \lambda)}{P(O | \lambda)}$$


using Bayes' theorem

This numerator is equal to the product of the forward and backward variables for state  $i$  at time  $t$ , which we previously described.

$$\gamma_t(i) = P(q_t = i | O, \lambda) = \frac{\alpha_t(i)\beta_t(i)}{P(O | \lambda)}$$

- The forward variable  $\alpha_t(i)$  takes into account the partial observation sequence  $\{o_1 o_2 \dots o_t\}$  ending at state  $i$  at time  $t$
- The backward variable  $\beta_t(i)$  takes into account the remainder of the observation  $\{o_{t+1} o_{t+2} \dots o_T\}$  given state  $i$  at time  $t$ .

# Learning: Baum-Welch Algorithm

Now, let's define  $\xi_t(i, j)$  as the probability of having a state transition from  $i$  to  $j$  at time  $t$ , given an observation sequence.

$$\xi_t(i, j) = P(q_t = i, q_{t+1} = j | O, \lambda) = \frac{P(O, q_t = i, q_{t+1} = j | \lambda)}{P(O | \lambda)}$$

Again, using Bayes' theorem

The numerator is similar to what we found for  $\gamma_t(i)$  and can again be simplified:

$$\xi_t(i, j) = \frac{\alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)}{P(O | \lambda)}$$

- The forward variable  $\alpha_t(i)$  takes into account the partial observation sequence  $\{o_1 o_2 \dots o_t\}$  ending at state  $i$  at time  $t$
- The backward variable  $\beta_{t+1}(j)$  takes into account the remainder of the observation  $\{o_{t+1} o_{t+2} \dots o_T\}$  given state  $j$  at time  $t + 1$

# Learning: Baum-Welch Algorithm

Now we can explore the re-estimation of the HMM model.

- A re-estimated model,  $\lambda^*(A, B, \pi)$  can be derived to maximize  $P(O|\lambda)$  such that  $P(O|\lambda^*) > P(O|\lambda)$
- This process is repeated iteratively, each time updating the model based on the previous estimates
- Unfortunately, the likelihood function for an HMM is quite complex and has many local maxima to which the re-estimation may converge.

We'll explain concepts without getting too deep into these, but be aware

The re-estimated initial state probabilities,  $\pi_i^*$ , are just:

$\pi_i^*$  = Expected number of times in state  $i$  at time ( $t = 1$ )

$$\pi_i^* = \gamma_1(i)$$

# Learning: Baum-Welch Algorithm

The re-estimated transition probabilities,  $a_{ij}^*$ , are given by:

$$a_{ij}^* = \frac{\text{expected number of transitions from state } i \text{ to } j}{\text{expected number of transitions from state } i} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \gamma_t(i)}$$

- Recall that  $\xi_t(i, j)$  is the probability of going from state  $i$  to state  $j$  at time  $t$ .
- So, the numerator (the sum of  $\xi_t(i, j)$  over all time) is the expected number of transitions from state  $i$  to state  $j$ .
- Recall that  $\gamma_t(i)$  is the probability of being in state  $i$  at time  $t$ .
- Then, the denominator (the sum of  $\gamma_t(i)$  over all time) is the expected number of transitions from state  $i$ .

Note: The summations don't include  $T$  because there are no transitions from time  $T$

# Learning: Baum-Welch Algorithm

The re-estimated **transition probabilities**,  $a_{ij}^*$ , are given by:

$$a_{ij}^* = \frac{\text{expected number of transitions from state } i \text{ to } j}{\text{expected number of transitions from state } i} = \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \gamma_t(i)}$$

Substituting in the formulas for  $\xi_t(i, j)$  and  $\gamma_t(i)$ , we get:

$$a_{ij}^* = \frac{\sum_{t=1}^{T-1} \frac{\alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)}{P(O|\lambda)}}{\sum_{t=1}^{T-1} \frac{\alpha_t(i) \beta_t(i)}{P(O|\lambda)}} = \frac{\sum_{t=1}^{T-1} \alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)}{\sum_{t=1}^{T-1} \alpha_t(i) \beta_t(i)}$$

So, the result is simply composed of:

- The forward result
- The backward result
- The previous transition probabilities
- The previous observation probabilities



# Learning: Baum-Welch Algorithm

The re-estimated **observation probabilities**,  $b_j^*(k)$ , are given by:

$$b_j^*(k) = \frac{\text{expected number of times observing } o_k \text{ when in state } j}{\text{expected number of times in state } j} = \frac{\sum_{t=1, o_k}^T \gamma_t(j)}{\sum_{t=1}^T \gamma_t(j)}$$

- The denominator (the sum of  $\gamma_t(j)$  over all time) is the same as before (the expected number of transitions from state  $j$ ), except that it now includes the full 1:T
- The numerator is similar to the denominator, except that we only include terms for which we are in state  $j$  AND observing symbol  $o_k$
- If we define  $\delta(o_t, o_k) = \begin{cases} 1 & \text{if } o_t = o_k \\ 0 & \text{otherwise} \end{cases}$ , and substitute the formula derived for  $\gamma_t(i)$ ,

$$b_j^*(k) = \frac{\sum_{t=1}^T \frac{\alpha_t(i)\beta_t(i)}{P(O|\lambda)} \delta(o_t, o_k)}{\sum_{t=1}^T \frac{\alpha_t(i)\beta_t(i)}{P(O|\lambda)}} = \frac{\sum_{t=1}^T \alpha_t(i)\beta_t(i)\delta(o_t, o_k)}{\sum_{t=1}^T \alpha_t(i)\beta_t(i)}$$

# Learning: Baum-Welch Algorithm

Using these formulas, we can iterate over the re-estimation until the results converge, such that  $P(O|\lambda^*) \approx P(O|\lambda)$

$$\pi_i^* = \gamma_1(i)$$

$$a_{ij}^* = \frac{\sum_{t=1}^{T-1} \alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)}{\sum_{t=1}^{T-1} \alpha_t(i) \beta_t(i)} \quad b_j^*(k) = \frac{\sum_{t=1}^T \alpha_t(i) \beta_t(i) \delta(o_t, o_k)}{\sum_{t=1}^T \alpha_t(i) \beta_t(i)}$$

Note that these derivations have assumed that the observations were discrete

- We could count the number of times each observation occurred for each state



# Continuous Observation Densities

Note that these derivations have assumed that the observations were discrete

- We could count the number of times each observation occurred for each state

In practice, this is often not true, requiring the consideration of continuous observation densities

- $\pi$  and  $A$  stay the same, but now we estimate the probability density function (pdf) for each state, instead of the (discrete) probability mass function,  $b_j(k)$ .



# Continuous Observation Densities

Whereas re-estimating the pmf,  $b_j^*(k)$  was achieved using

$$b_j^*(k) = \frac{\sum_{t=1}^T \alpha_t(i) \beta_t(i) \delta(o_t, o_k)}{\sum_{t=1}^T \alpha_t(i) \beta_t(i)}$$

Estimation of the pdf is most often achieved using [Gaussian Mixture Models](#)

- We describe the pdf as a weighted sum of  $M$  Gaussian distributions,  $\eta$

$$b_j(\bar{o}) = \sum_{k=1}^M c_{jk} \eta(\bar{o}, \bar{\mu}_{jk}, \bar{U}_{jk}) \quad 1 \leq j \leq N$$

- Here,  $\bar{\mu}_{jk}$  is the mean vector, and  $\bar{U}_{jk}$  is the covariance matrix for the  $k^{\text{th}}$  mixture component in state  $j$ .
- The coefficients  $c_{jk}$  are the mixture coefficients, which weight the contribution of the Gaussians, subject to:

$$\sum_{k=1}^M c_{jk} = 1 \quad c_{jk} \geq 0 \quad 1 \leq j \leq N, 1 \leq k \leq M$$

# Continuous Observation Densities

The training of a continuous observation HMM is similar to before, but requires the additional re-estimation of these  $c_{jk}$ ,  $\bar{\mu}_{jk}$ , and  $\bar{U}_{jk}$  parameters.

As before, we need to defined a helpful variable.

- $\gamma_t(j, k)$  is the probability of being in state  $j$  at time  $t$  with the  $k^{\text{th}}$  mixture component accounting for the observation vector  $\bar{o}_t$

$$\gamma_t(j, k) = \left[ \frac{\alpha_t(j)\beta_t(j)}{\sum_{j=1}^N \alpha_t(j)\beta_t(j)} \right] \left[ \frac{c_{jk}\eta(\bar{o}_t, \bar{\mu}_{jk}, \bar{U}_{jk})}{\sum_{m=1}^M c_{jm}\eta(\bar{o}_t, \bar{\mu}_{jm}, \bar{U}_{jm})} \right]$$

- The first term is exactly the same as the discrete HMM equivalent,  $\gamma_t(j)$ .
- The second term is the ratio between the probability that the  $k^{\text{th}}$  mixture of state  $j$  can account for the observation vector  $\bar{o}_t$  and the probability that the observation vector  $\bar{o}_t$  can be accounted for in state  $j$ , in general.

# Continuous Observation Densities

The mixture gains,  $c_{jk}$ , can then be re-estimated as


$$c_{jk} = \frac{\sum_{t=1}^T \gamma_t(j, k)}{\sum_{t=1}^T \sum_{k=1}^M \gamma_t(j, k)}$$

- which is the ratio between the expected number of times the system is in state  $j$ , using the  $k^{\text{th}}$  mixture component, and the total number of times the system is in state  $j$ .

The means and covariances of the Gaussian mixtures can be re-estimated using

$$\bar{\mu}_{jk} = \frac{\sum_{t=1}^T \gamma_t(j, k) \cdot \bar{o}_t}{\sum_{t=1}^T \gamma_t(j, k)} \quad \bar{U}_{jk} = \frac{\sum_{t=1}^T \gamma_t(j, k) \cdot (\bar{o}_t - \bar{\mu}_{ij})(\bar{o}_t - \bar{\mu}_{ij})'}{\sum_{t=1}^T \gamma_t(j, k)}$$

- where the numerators are the sum of the probabilities of being in state  $j$ , using the  $k^{\text{th}}$  mixture, weighted by the observation vector  $\bar{o}_t$  (for the mean) or the covariance.



# Hidden Markov Models

Using priori knowledge about the problem, we can constrain HMMs for better performance or efficiency. We can:


- constrain models to be left-right model (state index is non-decreasing)

$$a_{ij} = 0 \quad j < i$$


- limit how far 'ahead' a model can jump

$$a_{ij} = 0 \quad j > i + \Delta i$$


- Force the model to start in state 1, or end in state N



left-to-right model



parallel path left-to-right model



ergodic model

$$\begin{pmatrix} a_{11} & a_{12} & 0 \\ 0 & a_{22} & a_{23} \\ 0 & 0 & a_{33} \end{pmatrix}$$

$$\begin{pmatrix} a_{11} & a_{12} & a_{13} & 0 & 0 \\ 0 & a_{22} & a_{23} & a_{24} & 0 \\ 0 & 0 & a_{33} & a_{34} & a_{35} \\ 0 & 0 & 0 & a_{44} & a_{45} \\ 0 & 0 & 0 & 0 & a_{55} \end{pmatrix}$$

$$\begin{pmatrix} a_{11} & a_{12} & a_{13} & a_{14} & a_{15} \\ a_{21} & a_{22} & a_{23} & a_{24} & a_{25} \\ a_{31} & a_{32} & a_{33} & a_{34} & a_{35} \\ a_{41} & a_{42} & a_{43} & a_{44} & a_{45} \\ a_{51} & a_{52} & a_{53} & a_{54} & a_{55} \end{pmatrix}$$

# Hierarchical Hidden Markov Models

Complex HMMs can be designed to be more efficient by carefully constraining the transitions matrix,  $A$

- However, they can sometimes be hard to interpret directly

Several extensions have therefore been proposed, particularly to embed or emphasize known structure.

A particularly interesting (and more common) variant is the hierarchical HMM

- As the name implies, hHMMs combines multiple HMMs in a hierarchical (or tree-like) arrangement
- Each state of the original (top) model is comprised of its own HMM
- In this way, the observations associated with a given state in the upper level(s) are informed by the state sequence through the associated 'leaf' HMM.

As an example, consider the English language

- Sentences are formed from sequences of words
- Words are formed from sequences of syllables
- Syllables are formed from sequences of phonemes




# Hierarchical Hidden Markov Models

A hierarchical HMM could then be designed where the top level models the flow of words through a sentence

- For each word, a new hierarchical HMM could model the flow of syllables within that word
- And a final HMM could model the flow of phonemes through those syllables

Within each level of a hierarchical HMM, there is a sequence of states



- An output is only returned to the next level up once the sequence has been completed
- That is, the first and middle states of each HMM are 'non-producing', and only inform the observation that is generated by the last state.



# Hierarchical Hidden Markov Models

Recent applications of hHMMs include

- Speech
- Gesture and human activity recognition
- Gait
- Genetic sequencing



Fine, Shai, Singer, Yoram, Tishby, Naftali, The Hierarchical Hidden Markov Model: Analysis and Applications, Machine Learning, 32, 41–62 (1998)

Martindale, Christine & Hoenig, Florian & Strohrmann, Christina & Eskofier, Bjoern. (2017). Smart Annotation of Cyclic Data Using Hierarchical Hidden Markov Models. Sensors. 17. 2328. 10.3390/s17102328.

# Summary

In this section, we continued with another common state-space model; the Hidden Markov Model

- A sequence of latent states, which each generate some observable output  
The observations can be deterministic, discrete, or continuous (based on a pdf)
- State progression is governed by the probabilities in a state transition matrix  
Observations depend only on the current state
- HMMs can be used for time series modeling, or for classification of temporal sequences  
We focused on the classification task
- Three main “problems”  
Evaluation (forward, backward), Decoding (Viterbi), Learning (Baum-Welch)
- The models can be constrained for more efficient learning or to better represent known structure/semantics  
Left-right vs ergodic models, hierarchical HMMs

In the next section, we’ll jump again and look at deep learning models

- Regression & Gradient Descent
- RNN, LSTM, maybe CNN (as time allows)

# Q&A

